A. Generate held-out features: repeat with five disjoint-group rotations

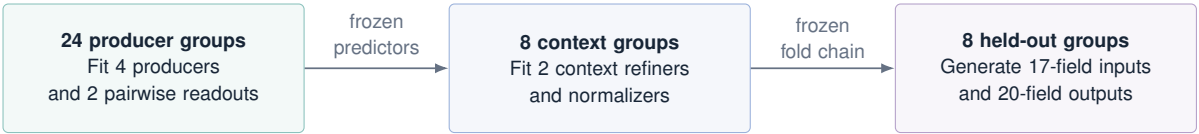

Arrows transfer fitted predictors, not training sequences. Held-out targets enter only after feature generation.

B. Fit the deployment components from three distinct training inputs

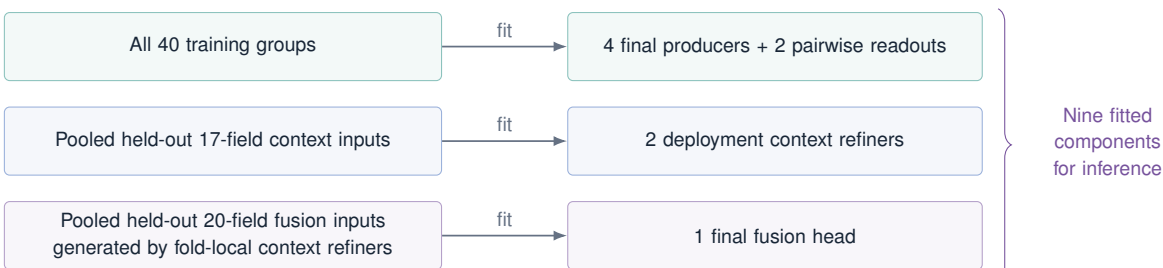

Deployment context refiners do not generate the fusion head's training inputs.